

# Just in Time Clouds


Rostand Costa<sup>1,2</sup>, **Francisco Brasileiro**<sup>1</sup>  
Guido Lemos Filho<sup>2</sup>, Dênio Mariz Sousa<sup>2,3</sup>

<sup>1</sup>Universidade Federal de Campina Grande

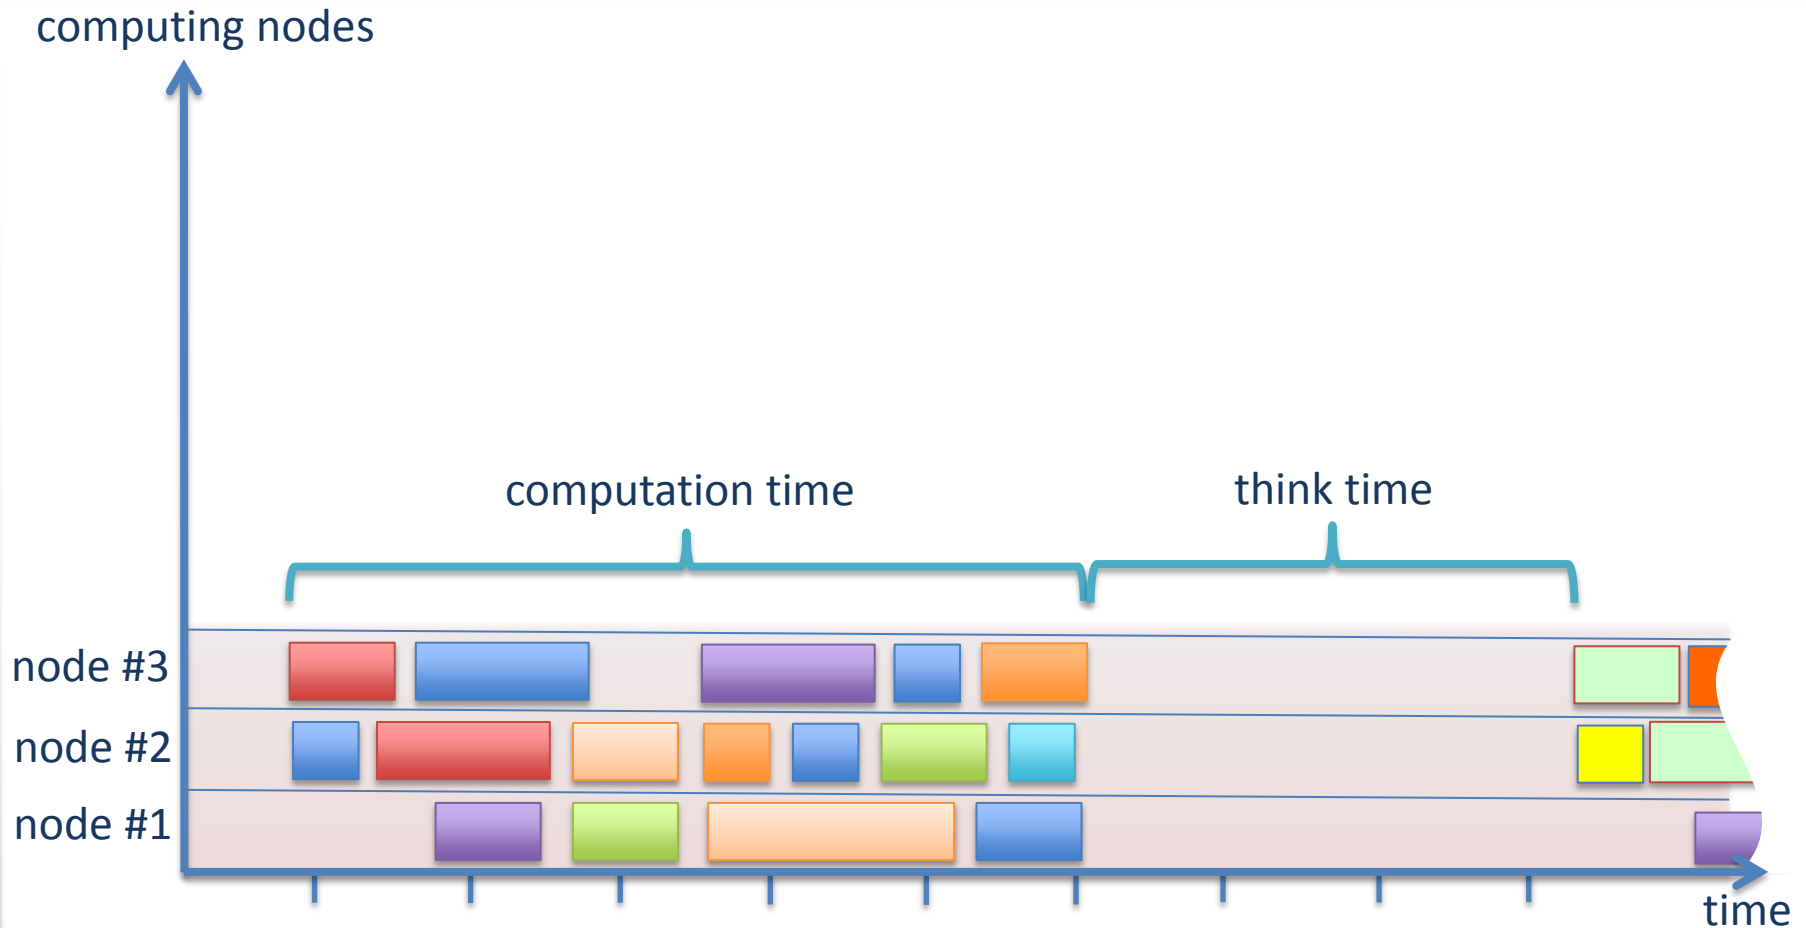
<sup>2</sup>Universidade Federal da Paraíba

<sup>3</sup>Instituto Federal de Educação, Ciência e Tecnologia da Paraíba

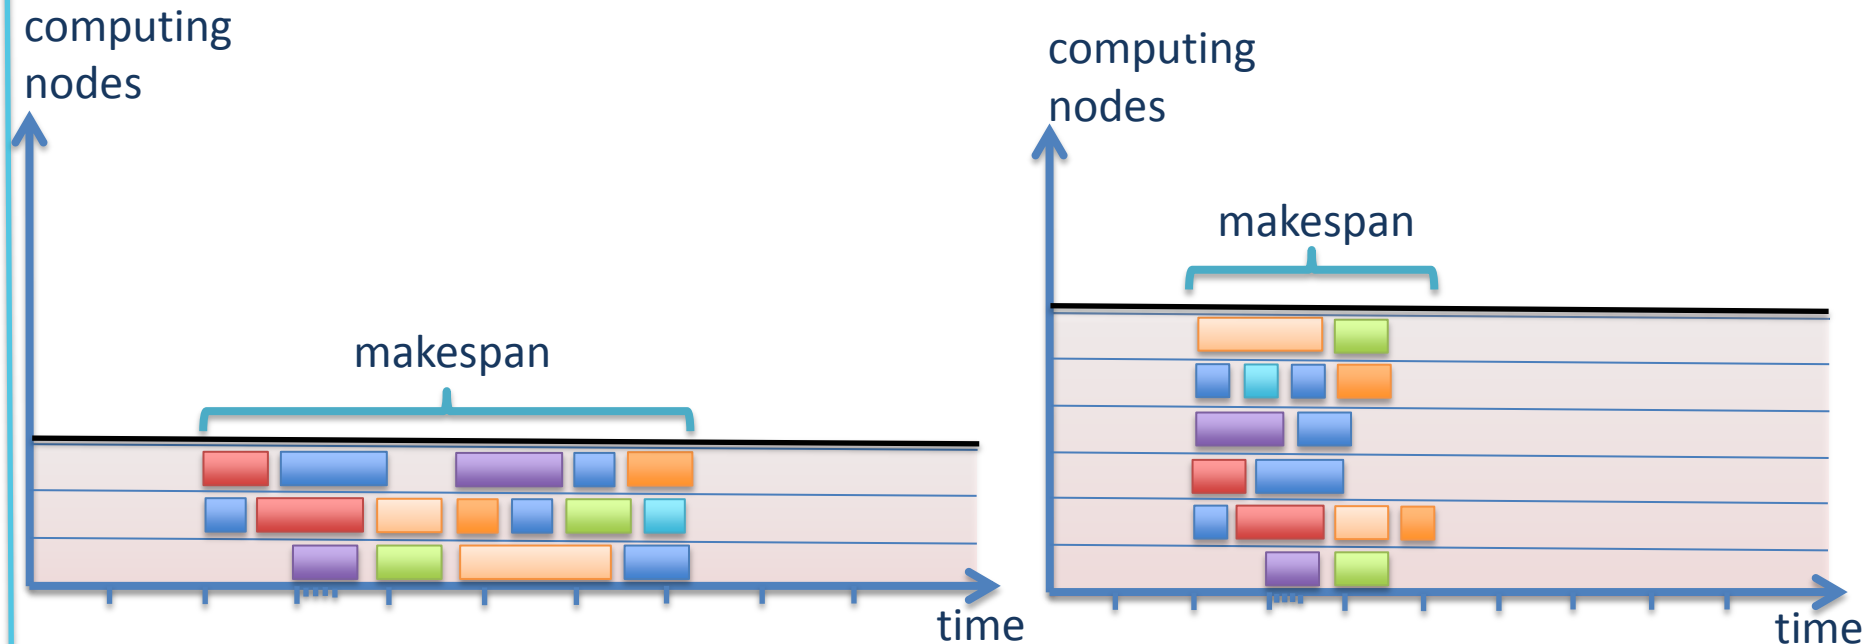
# Bag-of-task applications



# Typical workload of BoT applications




# Trade-off performance vs idleness level




# Running BoT on cloud computing IaaS providers

computing nodes


Pay-as-you-go and elasticity properties eliminates the trade-off



# Cost associativity of cloud providers



Same charge for any two requests  $s_i$  and  $s_j$ ,  
with  $\rho_i \times \delta_i = \rho_j \times \delta_j$



# How elastic can one be in practice?

**20 on-demand instances**

If you wish to run more than 20 On-Demand or Reserved Instances or 100 Spot Instances, create more than 100 EBS volumes, need more than 5 Elastic IP addresses or 5 Elastic Load Balancers, or need to send large quantities of email from your EC2 account, please complete the [Amazon EC2 instance request form](#), [Amazon EBS volume request form](#), [Elastic IP request form](#), [Elastic Load Balancers](#), or the [Email request form](#) respectively and your request will be considered.

**100 “spot” instances**

## Request to Increase Amazon EC2 Instance Limit

Please complete the form below to request an increase to your Amazon EC2 instance limit. An asterisk (\*) indicates required information:

First Name\*

Last Name\*

Email Address\*

Telephone\*

Company Name\*

Company Size\*

Country\*

State/Province\*

Postal Code\*

AWS Account ID

Email Address Associated with AWS Account\*

Requested New Instance Limit (Number)\*

EC2 Regions\*

Operating System\*

Primary Instance Type\*

Frequency of Usage\*

Use Case Description\*

Please include as much detail as possible, including duration of use.

# This led us to our first research question ...

- Why public cloud providers impose limits that restrict the usefulness of their services for clients with BoT applications?
  - How this limit affects the minimum capacity required by the provider?
  - How the rate of the infrastructure's idleness increases with the increase of this limit?
  - How extremely elastic usage patterns impact the cost for the provider to run its infrastructure?



# Methodology

- Our approach is based on the use of simulation
- We defined a simplified model for IaaS providers ...  
... and, in the absence of real workload traces, an appropriate synthetic workload generator for the proposed model
- We performed an experimental design to identify the random variables that had a major impact on the response variables of the model ...  
... and, in the absence of data from cloud providers, we performed a parameter sweep over the main factors identified

# Model sketch

- We follow a business-driven approach:
  - The provider's profit over a period of time  $\Delta T$  is given by:
    - the revenue of the capacity sold during  $\Delta T$   
MINUS
    - the cost of maintaining the infrastructure during  $\Delta T$   
MINUS
    - the extra cost incurred only when its capacity is actually used during  $\Delta T$   
MINUS
    - the cost of any violations incurred during  $\Delta T$  (e.g. not being able to serve a client's request)

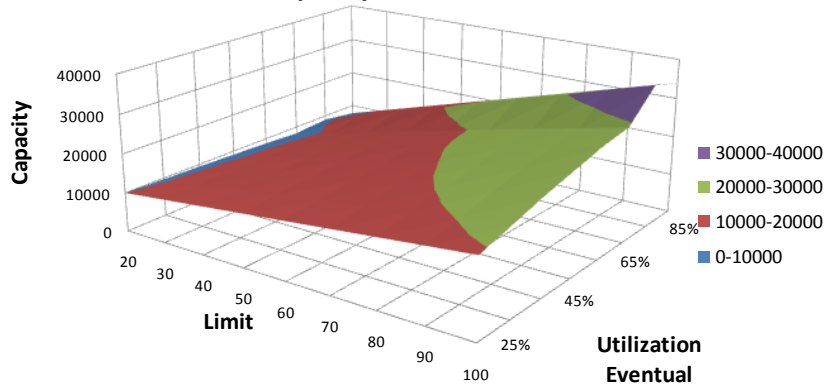
# Synthetic workload: user-based model

→ *Hierarchical Generative Model*


# Utilization profiles of the synthetic workload population

# Experiment #1: results


Minimum Capacity for 10% of BoT Users



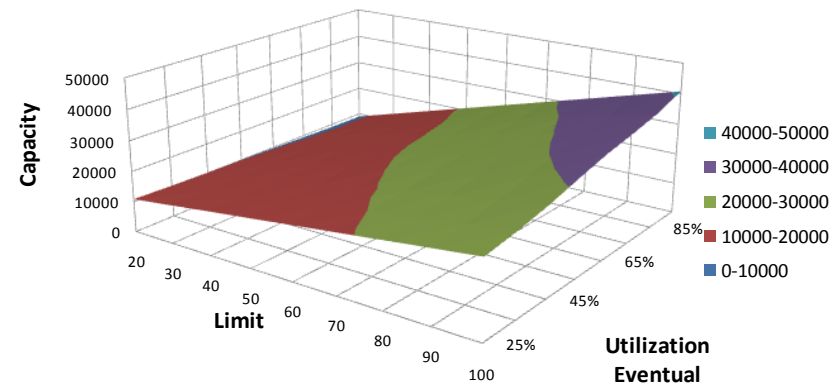
Minimum Capacity for 15% of BoT Users



Minimum Capacity for 20% of BoT Users




Minimum Capacity for 25% of BoT Users




# Experiment #2: results


## Idleness with 10% of BoT Users



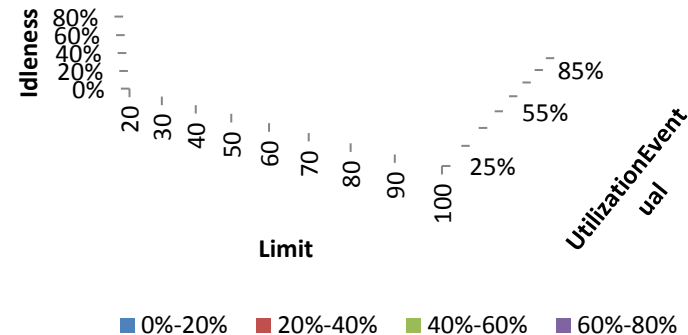
## Idleness with 15% of BoT Users



## Idleness with 20% of BoT Users

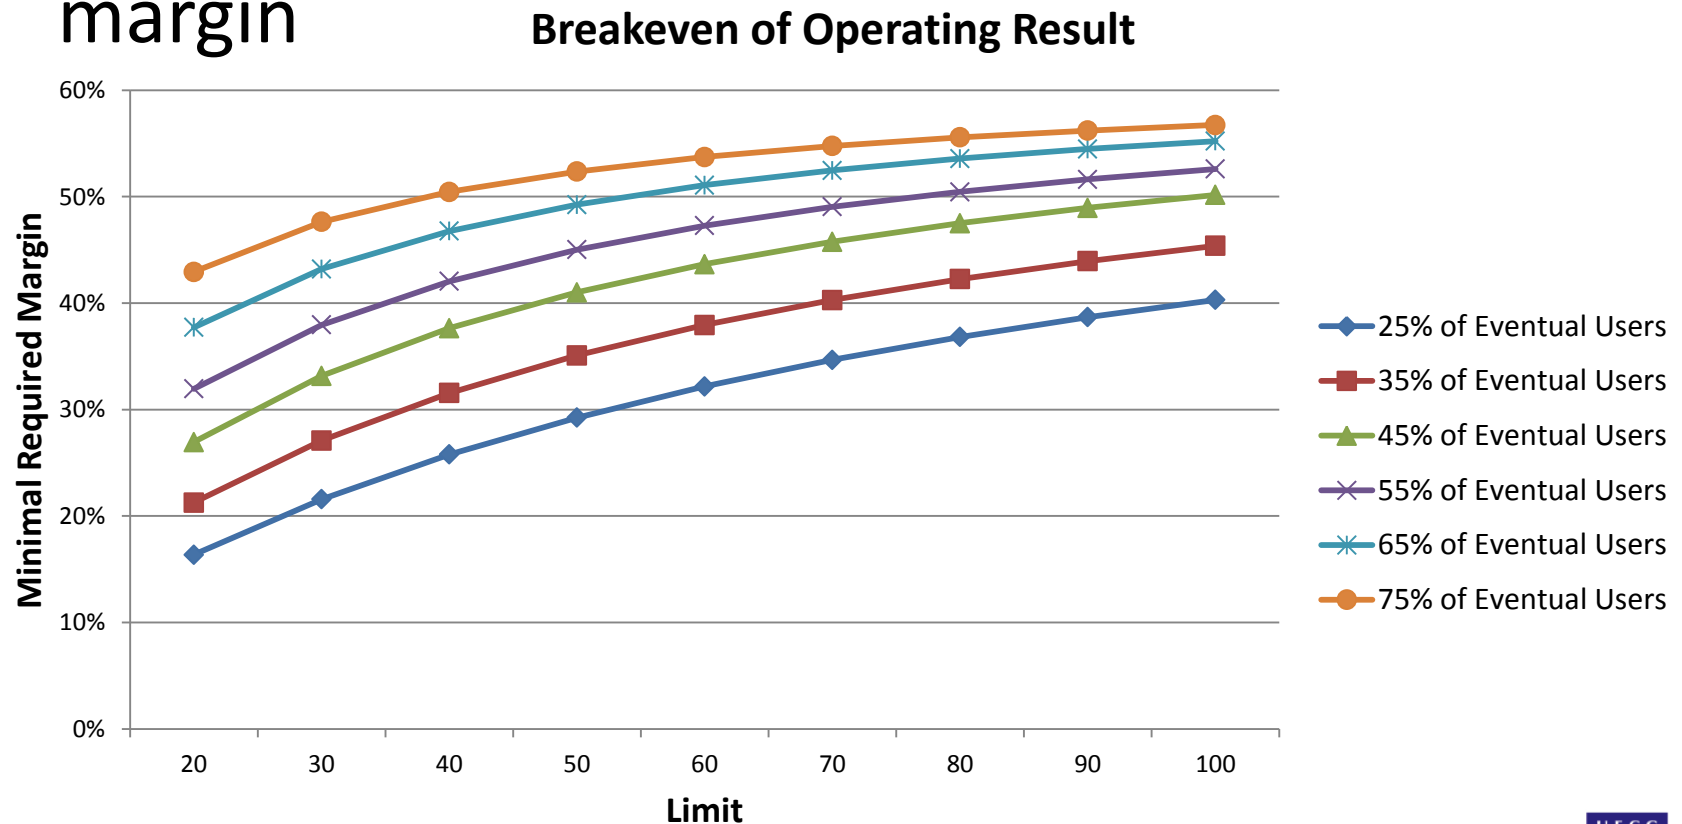


## Idleness with 25% of BoT Users



# Experiment #3: results

- Impact of the limit on the minimal required margin



# Partial conclusions

- The limit value has a **relevant impact** on the cost of running the infrastructure
  - Availability cost
- Users with very intensive and eventual use (BoT) **press the minimum capacity** and increase the system idleness
- Maintained the same profile of the population and the limit value, the **system dynamics is not dependent** on the number of users
  - Scaling does not eliminate the problem



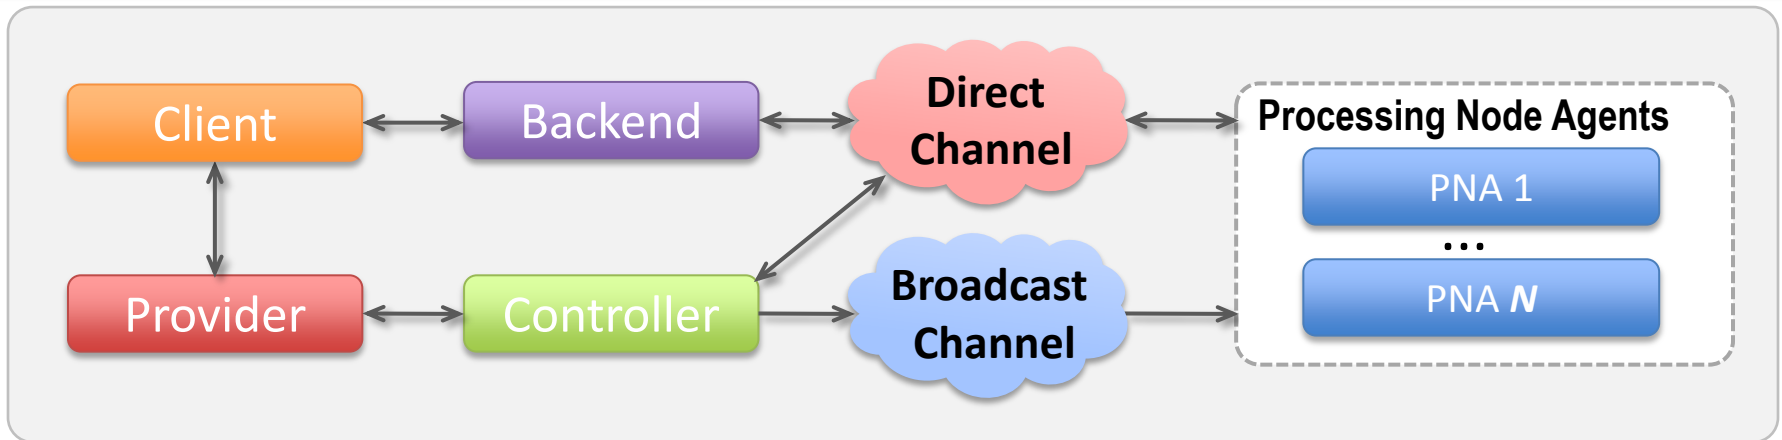
# How can we appropriately serve BoT users in an IaaS setting?

- We have been investigating alternative approaches to build highly elastic IaaS providers
  - Federate resources **owned by third parties** that are not able to provide themselves an IaaS infrastructure over the eventual spare capacity existent
    - Availability cost may be irrelevant because it has already been amortized by the main business supported by the resources
      - Small datacenters, P2P grids, unconventional devices (STBs in DTV systems, smartphones, etc.)
- Cloud infrastructure is assembled **“Just in Time”** by a broker

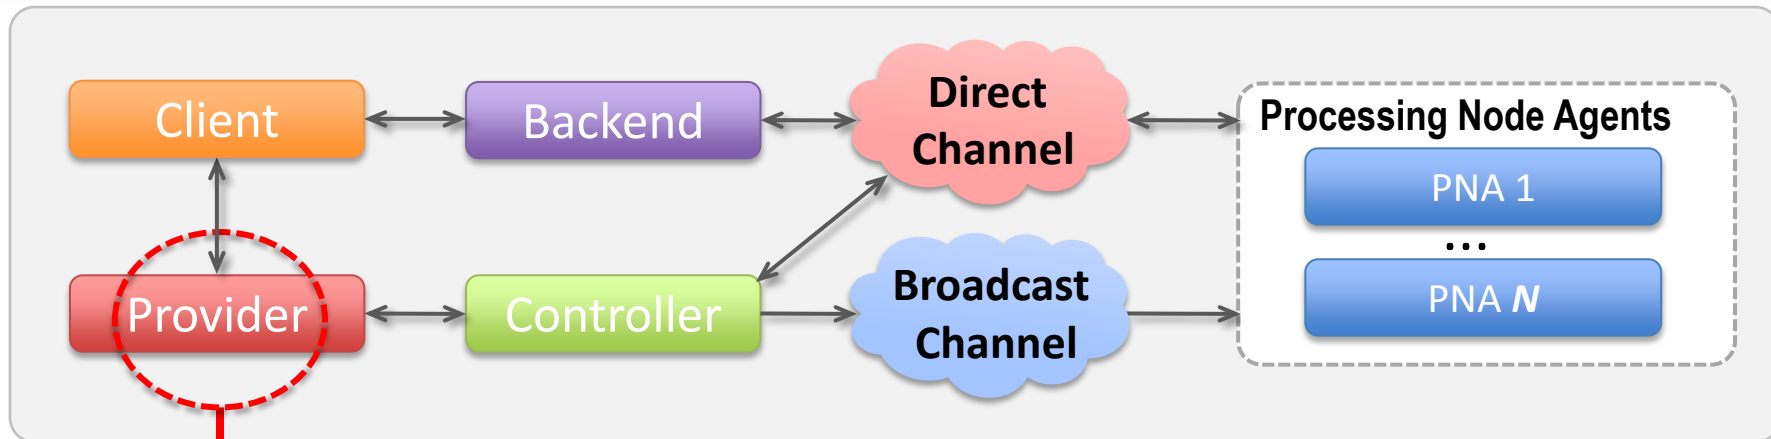
# On-Demand Distributed Computing Infrastructure

- OddCI considers a special category of devices which may be organized as a **broadcast network**
  - Mobile phones, Digital TV receivers, Cable TV receivers
  - Devices connected to the Internet with reasonably powerful processors
  - Broadcast network can access **simultaneously** all the devices which can be coordinated to run some task

# OddCI Architecture

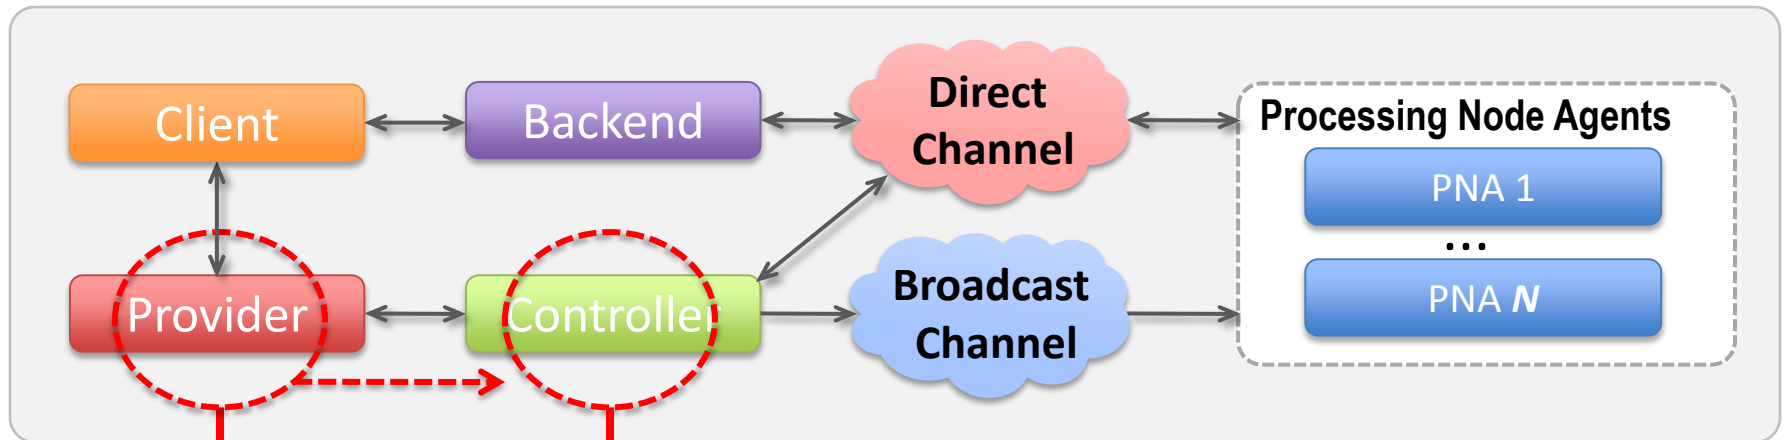


# OddCI Architecture: operation



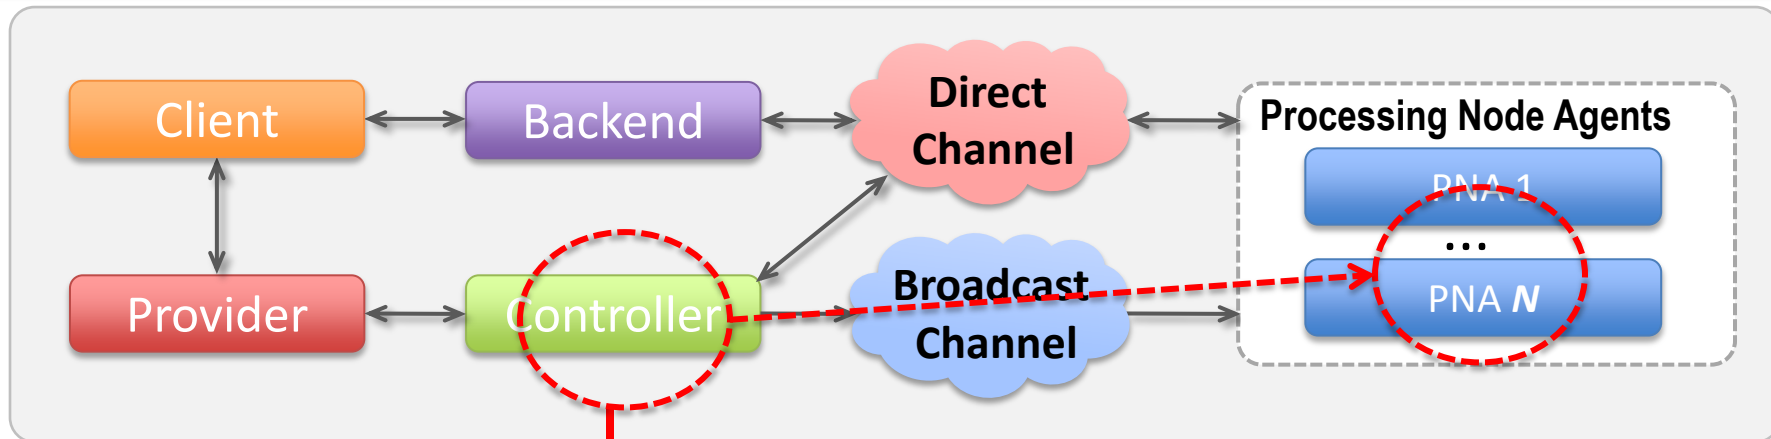
- Client submits a “processing request” to the **provider**
  - DCI instance size (number of processing nodes)
  - Application image, common data
  - Node requirements

# OddCI Architecture: operation



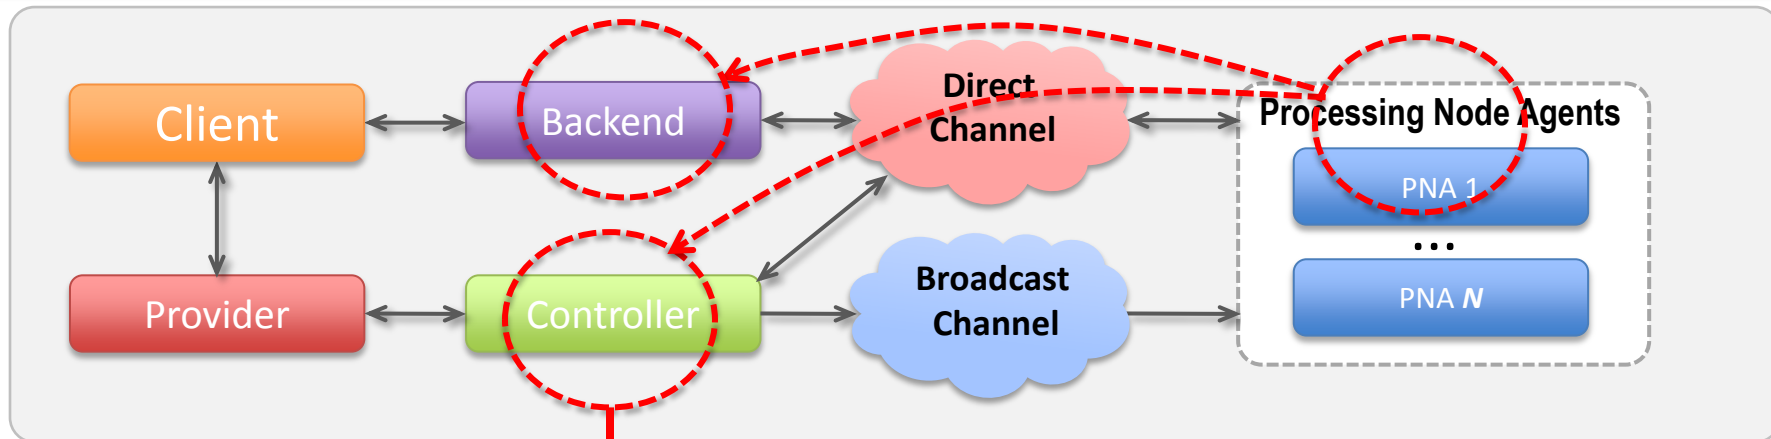
- **Provider** evaluates the client's request
  - authenticates the client
  - checks availability
  - keeps control information
  - commands the **Controller** for creating the required OddCI instance

# OddCI Architecture: operation



- **Controller** triggers a wakeup message to **PNA**s through the **broadcast** channel
  - Wakeup message carry the PNA software and user application
- All **PNA** receive messages virtually at the same time
- **Controller** also sends other control messages (e.g. dismantle instances, leave, ...)

# OddCI Architecture: operation



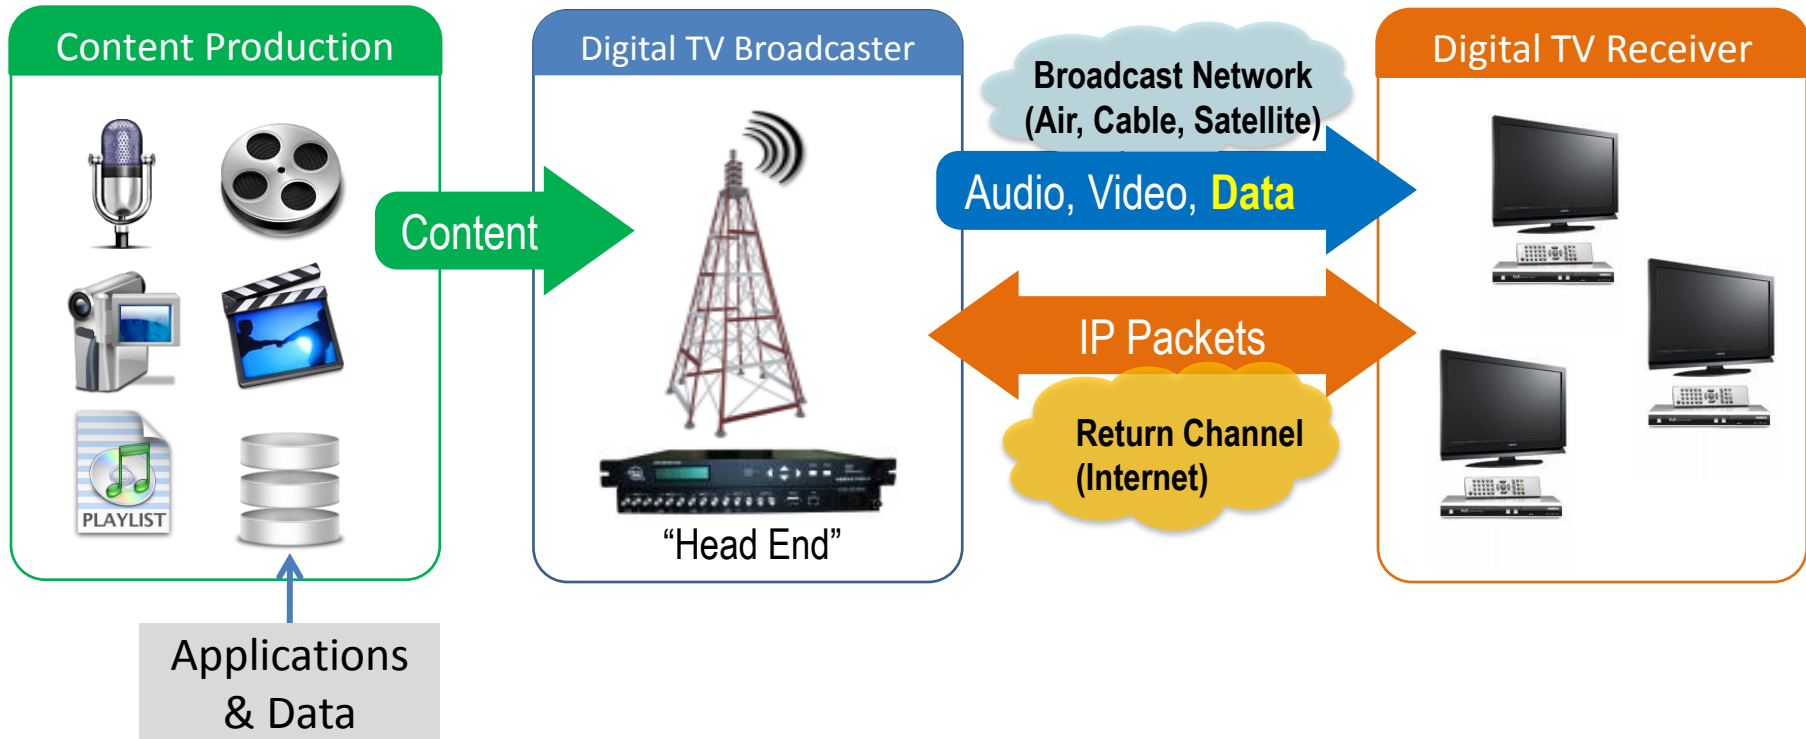
- **PNA** loads application image for execution in a DVE (Dynamic Virtual Environment)
- **Controller** monitors active **PNA**
- Use of the **Direct channel**
  - **Application** can interact with the **Backend** for requesting specific input data or send results
  - **PNA** sends status messages frequently to the **Controller**

# OddCI-Ginga: OddCI over a Digital TV Network

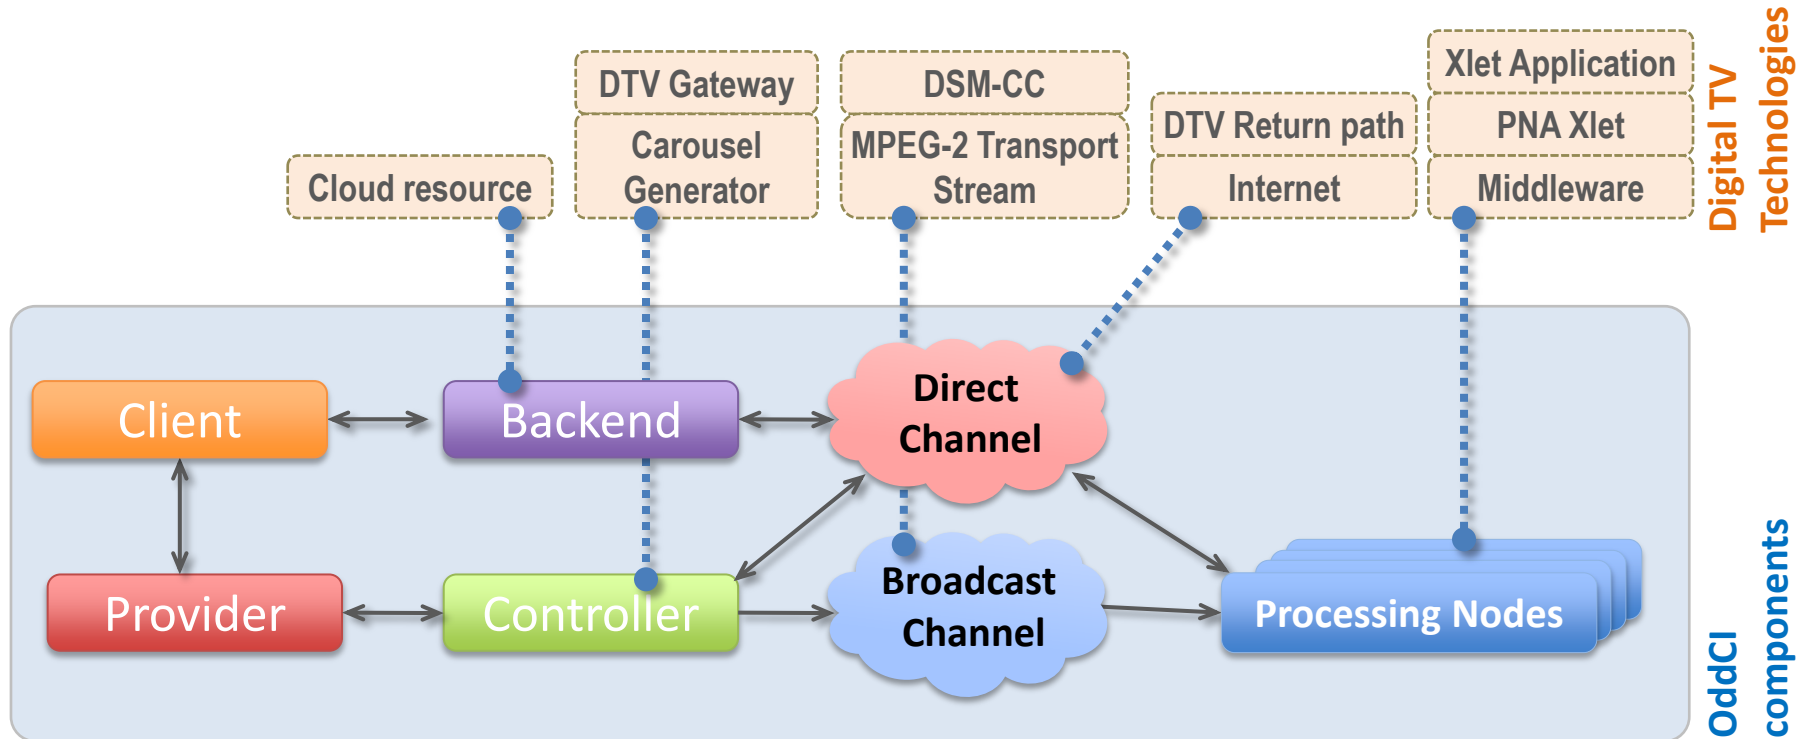
- Why DTV network?
  - Open technology, well-defined standards
  - Native transmission of data in broadcast
  - Fast expansion, being deployed in many countries
  - Potential for millions of devices
  - Powerful middleware
- And also ...
  - Feasibility for building a testbed
  - Previous experience of our group helping to develop Brazilian DTV middleware (SBTVD/**Ginga**)



# DTV Generic Model



# Implementing OddCI over DTV components



# Performance assessment: testbed

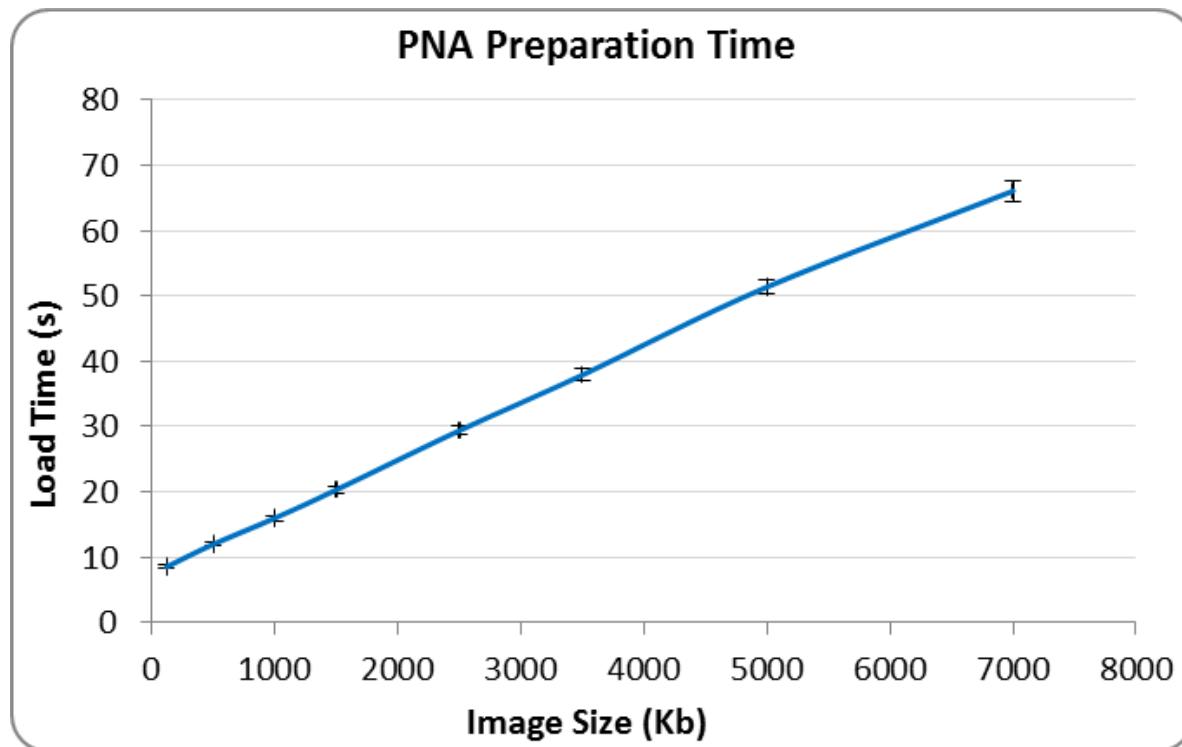
- TV Station
  - ISDB-T Digital Modulator (Brand/model: Linear ISMOD - Series ISCHIO)
  - Carousel Generator and Multiplexer (Brand/model: Domm Xstream)
  - Maximum data rate of data carousel set to 1Mbps
- Digital TV Receivers
  - Low-end: Proview XPS-1000, STi7001 processor 266MHz, 256MB RAM, Ethernet, STLinux
- Controller and Backend
  - Implemented as network services: Apache/Tomcat v6.0.33, HTTP for message exchange
  - Web framework Grails/Groovy scripts
  - MySQL v.5.1 for storing tasks and results in the backend.
- Provider
  - Web application for requesting the creation of instances and upload files to be sent to the PNA through the DTV data carousel

# Performance assessment: Benchmarking


- Bioinformatics application
  - We ported the NCBI Toolkit (blastall and blastcl3 programs) to the low-end DTV receiver used
- Reference machine (for comparison with DTV receiver)
  - Notebook Intel (R) Core i3-2310M 2.1GHz, 4 GB RAM, Fast Ethernet, Ubuntu 11.10 64-bits
  - Same applications were tested and compared

# Results: PNA load time (broadcast channel)

- Time for application load at PNA is linear with image (executable) size



# Results: DTV receiver Performance



Lower is better

When in “stand by”,  
DTV receiver is 65%  
faster than when  
“watching TV” mode

- Experiment Setup
  - Tests performed in a low-end DTV receiver (the **cheapest** in the Brazilian market)
- Remarks
  - Reference PC is ~23 times faster than DTV receiver in “in use” mode
  - Reference PC is ~12 times faster than DTV receiver in “standby” mode
  - Background applications does not interfere on the image/audio quality

# Concluding Remarks

- OddCI: a novel approach to DCI
  - Efficient setup, on-demand instantiation
  - Great potential to enable DCI for Extremely High-Throughput Computing
- OddCI can be instantiated over a DTV system
  - Less individual processing power, but huge pool size
    - Brazilian DTV expects ~100 million receivers by 2016
    - European DVB: >500 million receivers deployed
    - Chinese DTMB: ~100 million receivers (estimated)
- Many research challenges still opened

# Current Work

- Small scale deploy at the city of João Pessoa
  - 100 STBs to test interactivity of the SBTVD
- Peer-to-peer JiT Clouds
  - Combine the ideas of peer-to-peer grids (such as OurGrid) with the federation of clouds
  - Provide the same user experience that users have when use spot instances in Amazon Web Services
    - Different business model (market vs resource exchange)



# Thanks for your attention!

---

- Contact me at: [fubica@dsc.ufcg.edu.br](mailto:fubica@dsc.ufcg.edu.br)

Questions??